

Machine Learning LAB - Assessment 3

20BCE0083- Jeevan Yohan Varghese

1. K nearest neighbour

Write a program to implement k-Nearest Neighbour algorithm to classify the iris data set. Print both correct and wrong predictions.

```
In [21]: import numpy as np
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
```

```
In [22]: iris = datasets.load_iris()
X = iris.data
y = iris.target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
In [26]: def predict(X_train, y_train, x, k):
    #predicting class of a sample
    dist = np.sqrt(np.sum((X_train - x)**2, axis=1))
    nearest_indices = np.argsort(dist)[:k]
    nearest_labels = y_train[nearest_indices]
    return np.argmax(np.bincount(nearest_labels))

def predict_multiple(X_train, y_train, X_test, k):
    #utility function to predict multiple samples
    y_pred = np.array([predict(X_train, y_train, x, k) for x in X_test])
    return y_pred
```

```
In [27]: # k - no of neighbours
k = 10

y_pred = predict_multiple(X_train, y_train, X_test, k)

# accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.9

```
In [28]: #predicting all
for i in range(len(y_pred)):
    if y_pred[i] == y_test[i]:
        print("Sample", i, ": correct prediction as class", y_pred[i])
    else:
        print("Sample", i, "incorrect prediction as class", y_pred[i], "actu
```

```

Sample 0 : correct prediction as class 1
Sample 1 : correct prediction as class 1
Sample 2 : correct prediction as class 2
Sample 3 : correct prediction as class 1
Sample 4 : correct prediction as class 1
Sample 5 : correct prediction as class 0
Sample 6 : correct prediction as class 0
Sample 7 : correct prediction as class 0
Sample 8 : correct prediction as class 1
Sample 9 incorrect prediction as class 1 actual class : 2
Sample 10 : correct prediction as class 0
Sample 11 : correct prediction as class 2
Sample 12 : correct prediction as class 1
Sample 13 : correct prediction as class 2
Sample 14 : correct prediction as class 2
Sample 15 : correct prediction as class 1
Sample 16 : correct prediction as class 2
Sample 17 incorrect prediction as class 1 actual class : 2
Sample 18 : correct prediction as class 1
Sample 19 : correct prediction as class 0
Sample 20 : correct prediction as class 0
Sample 21 : correct prediction as class 2
Sample 22 : correct prediction as class 0
Sample 23 : correct prediction as class 1
Sample 24 : correct prediction as class 1
Sample 25 : correct prediction as class 0
Sample 26 : correct prediction as class 1
Sample 27 : correct prediction as class 1
Sample 28 : correct prediction as class 0
Sample 29 incorrect prediction as class 2 actual class : 1

```

2. SVM

Train SVM classifier using sklearn digits dataset(i.e from sklearn datasets import load_digits)

```

In [11]: from sklearn import svm
         from sklearn.datasets import load_digits
         from sklearn.model_selection import train_test_split
         from sklearn.metrics import accuracy_score

```

```

In [12]: digits_data = load_digits()

         X_train, X_test, y_train, y_test = train_test_split(digits_data.data, digits_

```

```

In [13]: rbf = svm.SVC(kernel='rbf')
         rbf.fit(X_train, y_train)
         y_pred_rbf = rbf.predict(X_test)
         rbf_accuracy = accuracy_score(y_test, y_pred_rbf)
         print("RBF kernel accuracy:", rbf_accuracy)

```

RBF kernel accuracy: 0.9916666666666667

```

In [14]: lr = svm.SVC(kernel='linear')

```

```
lr.fit(X_train, y_train)
y_pred_lr = lr.predict(X_test)
linear_accuracy = accuracy_score(y_test, y_pred_lr)
print("Linear kernel accuracy:", linear_accuracy)
```

Linear kernel accuracy: 0.9777777777777777

```
In [15]: rbf = svm.SVC(kernel='rbf',C=0.5,gamma=0.001)
rbf.fit(X_train, y_train)
y_pred_rbf = rbf.predict(X_test)
rbf_accuracy = accuracy_score(y_test, y_pred_rbf)
print("Tuned RBF kernel accuracy:", rbf_accuracy)
```

Tuned RBF kernel accuracy: 0.9944444444444445

```
In [16]: lr = svm.SVC(kernel='linear',C=1,gamma=0.001)
lr.fit(X_train, y_train)
y_pred_lr = lr.predict(X_test)
linear_accuracy = accuracy_score(y_test, y_pred_lr)
print("Tuned Linear kernel accuracy:", linear_accuracy)
```

Tuned Linear kernel accuracy: 0.9777777777777777

3. ANN

Build an Artificial Neural Network by implementing the Back propagation algorithm and test the same using appropriate data sets.

```
In [17]: import numpy as np
X = np.array([[2, 9], [1, 5], [3, 6]], dtype=float)
y = np.array([[92], [86], [89]], dtype=float)
X = X/np.amax(X,axis=0)
y = y/100
def sigmoid (x):
    return 1/(1 + np.exp(-x))
def derivatives_sigmoid(x):
    return x * (1 - x)
```

```
In [18]: epoch=5
lr=0.1
inputlayer_neurons = 2
hiddenlayer_neurons = 3
output_neurons = 1
wh=np.random.uniform(size=(inputlayer_neurons,hiddenlayer_neurons))
bh=np.random.uniform(size=(1,hiddenlayer_neurons))
wout=np.random.uniform(size=(hiddenlayer_neurons,output_neurons))
bout=np.random.uniform(size=(1,output_neurons))
```

```
In [19]: for i in range(epoch):

    hinp1=np.dot(X,wh)
    hinp=hinp1 + bh
    hlayer_act = sigmoid(hinp)
    outinp1=np.dot(hlayer_act,wout)
    outinp= outinp1+bout
```

```

output = sigmoid(outinp)

E0 = y-output
outgrad = derivatives_sigmoid(output)
d_output = E0 * outgrad
EH = d_output.dot(wout.T)
hiddengrad = derivatives_sigmoid(hlayer_act)
d_hiddenlayer = EH * hiddengrad

wout += hlayer_act.T.dot(d_output) *lr
wh += X.T.dot(d_hiddenlayer) *lr

print ("Epoch-", i+1, "Starting")
print("Input: \n" + str(X))
print("Actual: \n" + str(y))
print("Predicted Output: \n" ,output)
print ("Epoch-", i+1, "Ending\n")

```

Epoch- 1 Starting
Input:
[[0.66666667 1.]
 [0.33333333 0.55555556]
 [1. 0.66666667]]
Actual:
[[0.92]
 [0.86]
 [0.89]]
Predicted Output:
[[0.84845988]
 [0.83935393]
 [0.84840029]]
Epoch- 1 Ending

Epoch- 2 Starting
Input:
[[0.66666667 1.]
 [0.33333333 0.55555556]
 [1. 0.66666667]]
Actual:
[[0.92]
 [0.86]
 [0.89]]
Predicted Output:
[[0.84891601]
 [0.83978841]
 [0.84884992]]
Epoch- 2 Ending

Epoch- 3 Starting
Input:
[[0.66666667 1.]
 [0.33333333 0.55555556]
 [1. 0.66666667]]
Actual:
[[0.92]
 [0.86]
 [0.89]]
Predicted Output:
[[0.84936544]
 [0.84021661]
 [0.84929297]]
Epoch- 3 Ending

Epoch- 4 Starting
Input:
[[0.66666667 1.]
 [0.33333333 0.55555556]
 [1. 0.66666667]]
Actual:
[[0.92]
 [0.86]
 [0.89]]
Predicted Output:
[[0.84980831]

```
[0.84063868]
[0.84972956]]
Epoch- 4 Ending
```

Epoch- 5 Starting

```
Input:
[[0.66666667 1.          ]
 [0.33333333 0.55555556]
 [1.          0.66666667]]
```

Actual:

```
[[0.92]
 [0.86]
 [0.89]]
```

Predicted Output:

```
[[0.85024476]
 [0.84105475]
 [0.85015984]]
```

Epoch- 5 Ending

```
In [20]: print("Input: \n" + str(X))
print("Actual Output: \n" + str(y))
print("Predicted Output: \n" ,output)
```

```
Input:
[[0.66666667 1.          ]
 [0.33333333 0.55555556]
 [1.          0.66666667]]
```

Actual Output:

```
[[0.92]
 [0.86]
 [0.89]]
```

Predicted Output:

```
[[0.85024476]
 [0.84105475]
 [0.85015984]]
```

4. Bagging Ensembles including Bagged Decision Trees, Random Forest and Extra Trees

```
In [30]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.impute import SimpleImputer
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
%matplotlib inline
sns.set_style("whitegrid")
plt.style.use("fivethirtyeight")
```

```
In [31]: df = pd.read_csv("diabetes.csv")
df.head()
```

```
Out[31]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunc
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	2

```
In [32]: pd.set_option('display.float_format', '{:.2f}'.format)
df.describe()
```

```
Out[32]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigr
count	768.00	768.00	768.00	768.00	768.00	768.00	
mean	3.85	120.89	69.11	20.54	79.80	31.99	
std	3.37	31.97	19.36	15.95	115.24	7.88	
min	0.00	0.00	0.00	0.00	0.00	0.00	
25%	1.00	99.00	62.00	0.00	0.00	27.30	
50%	3.00	117.00	72.00	23.00	30.50	32.00	
75%	6.00	140.25	80.00	32.00	127.25	36.60	
max	17.00	199.00	122.00	99.00	846.00	67.10	

```
In [34]: categorical_val = []
continous_val = []
for column in df.columns:
    if len(df[column].unique()) <= 10:
        categorical_val.append(column)
    else:
        continous_val.append(column)
df.columns
```

```
Out[34]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
               'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
              dtype='object')
```

```
In [35]: feature_columns = [
    'Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
    'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age'
]

for column in feature_columns:
    print(f"{column} ==> Missing zeros : {len(df.loc[df[column] == 0])}")
```

```
Pregnancies ==> Missing zeros : 111
Glucose ==> Missing zeros : 5
BloodPressure ==> Missing zeros : 35
SkinThickness ==> Missing zeros : 227
Insulin ==> Missing zeros : 374
BMI ==> Missing zeros : 11
DiabetesPedigreeFunction ==> Missing zeros : 0
Age ==> Missing zeros : 0
```

```
In [36]: fill_values = SimpleImputer(missing_values=0, strategy="mean", copy=False)
df[feature_columns] = fill_values.fit_transform(df[feature_columns])

for column in feature_columns:
    print(f"{column} ==> Missing zeros : {len(df.loc[df[column] == 0])}")
```

```
Pregnancies ==> Missing zeros : 0
Glucose ==> Missing zeros : 0
BloodPressure ==> Missing zeros : 0
SkinThickness ==> Missing zeros : 0
Insulin ==> Missing zeros : 0
BMI ==> Missing zeros : 0
DiabetesPedigreeFunction ==> Missing zeros : 0
Age ==> Missing zeros : 0
```

```
In [37]: X = df[feature_columns]
y = df.Outcome

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, ran
```

```
In [44]: def evaluate(model, X_train, X_test, y_train, y_test):
    y_test_pred = model.predict(X_test)
    y_train_pred = model.predict(X_train)

    print("Training Results: \n")
    clf_report = pd.DataFrame(classification_report(y_train, y_train_pred, out
    print(f"Confusion Matrix:\n{confusion_matrix(y_train, y_train_pred)}")
    print(f"Accuracy:\n{accuracy_score(y_train, y_train_pred):.4f}")
    print(f"Classification Report:\n{clf_report}")

    print("Testing results: \n")
    clf_report = pd.DataFrame(classification_report(y_test, y_test_pred, out
    print(f"Confusion Matrix:\n{confusion_matrix(y_test, y_test_pred)}")
    print(f"Accuracy:\n{accuracy_score(y_test, y_test_pred):.4f}")
    print(f"Classification Report:\n{clf_report}")
```

```
In [39]: from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier
```

```
In [45]: tree = DecisionTreeClassifier()
bagging_clf = BaggingClassifier(estimator=tree, n_estimators=1500, random_st
bagging_clf.fit(X_train, y_train)

evaluate(bagging_clf, X_train, X_test, y_train, y_test)
```


Training Results:

Confusion Matrix:

```
[[349  0]
 [ 0 188]]
```

Accuracy:

1.0000

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	1.00	1.00	1.00	1.00	1.00
recall	1.00	1.00	1.00	1.00	1.00
f1-score	1.00	1.00	1.00	1.00	1.00
support	349.00	188.00	1.00	537.00	537.00

Testing results:

Confusion Matrix:

```
[[119 32]
 [ 24 56]]
```

Accuracy:

0.7576

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	0.83	0.64	0.76	0.73	0.76
recall	0.79	0.70	0.76	0.74	0.76
f1-score	0.81	0.67	0.76	0.74	0.76
support	151.00	80.00	0.76	231.00	231.00

```
In [46]: #Random Forest
from sklearn.ensemble import RandomForestClassifier
rf_clf = RandomForestClassifier(random_state=42, n_estimators=1000)
rf_clf.fit(X_train, y_train)
evaluate(rf_clf, X_train, X_test, y_train, y_test)
```

Training Results:

Confusion Matrix:

```
[[349  0]
 [ 0 188]]
```

Accuracy:

1.0000

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	1.00	1.00	1.00	1.00	1.00
recall	1.00	1.00	1.00	1.00	1.00
f1-score	1.00	1.00	1.00	1.00	1.00
support	349.00	188.00	1.00	537.00	537.00

Testing results:

Confusion Matrix:

```
[[123  28]
 [ 29  51]]
```

Accuracy:

0.7532

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	0.81	0.65	0.75	0.73	0.75
recall	0.81	0.64	0.75	0.73	0.75
f1-score	0.81	0.64	0.75	0.73	0.75
support	151.00	80.00	0.75	231.00	231.00

```
In [47]: #Extra Trees
from sklearn.ensemble import ExtraTreesClassifier
ex_tree_clf = ExtraTreesClassifier(n_estimators=1000, max_features=7, random
ex_tree_clf.fit(X_train, y_train)
evaluate(ex_tree_clf, X_train, X_test, y_train, y_test)
```

Training Results:

Confusion Matrix:

```
[[349  0]
 [ 0 188]]
```

Accuracy:

1.0000

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	1.00	1.00	1.00	1.00	1.00
recall	1.00	1.00	1.00	1.00	1.00
f1-score	1.00	1.00	1.00	1.00	1.00
support	349.00	188.00	1.00	537.00	537.00

Testing results:

Confusion Matrix:

```
[[124  27]
 [ 25  55]]
```

Accuracy:

0.7749

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	0.83	0.67	0.77	0.75	0.78
recall	0.82	0.69	0.77	0.75	0.77
f1-score	0.83	0.68	0.77	0.75	0.78
support	151.00	80.00	0.77	231.00	231.00

5. Boosting Ensembles including AdaBoost and Stochastic Gradient Boosting

```
In [48]: from sklearn.ensemble import AdaBoostClassifier
ada_boost_clf = AdaBoostClassifier(n_estimators=30)
ada_boost_clf.fit(X_train, y_train)
evaluate(ada_boost_clf, X_train, X_test, y_train, y_test)
```

Training Results:

Confusion Matrix:

```
[[310  39]
 [ 51 137]]
```

Accuracy:

0.8324

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	0.86	0.78	0.83	0.82	0.83
recall	0.89	0.73	0.83	0.81	0.83
f1-score	0.87	0.75	0.83	0.81	0.83
support	349.00	188.00	0.83	537.00	537.00

Testing results:

Confusion Matrix:

```
[[123  28]
 [ 27  53]]
```

Accuracy:

0.7619

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	0.82	0.65	0.76	0.74	0.76
recall	0.81	0.66	0.76	0.74	0.76
f1-score	0.82	0.66	0.76	0.74	0.76
support	151.00	80.00	0.76	231.00	231.00

```
In [49]: #Gradient Boosting
from sklearn.ensemble import GradientBoostingClassifier
grad_boost_clf = GradientBoostingClassifier(n_estimators=100, random_state=42)
grad_boost_clf.fit(X_train, y_train)
evaluate(grad_boost_clf, X_train, X_test, y_train, y_test)
```

Training Results:

Confusion Matrix:

```
[[342  7]
 [ 19 169]]
```

Accuracy:

0.9516

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	0.95	0.96	0.95	0.95	0.95
recall	0.98	0.90	0.95	0.94	0.95
f1-score	0.96	0.93	0.95	0.95	0.95
support	349.00	188.00	0.95	537.00	537.00

Testing results:

Confusion Matrix:

```
[[116  35]
 [ 26  54]]
```

Accuracy:

0.7359

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	0.82	0.61	0.74	0.71	0.74
recall	0.77	0.68	0.74	0.72	0.74
f1-score	0.79	0.64	0.74	0.72	0.74
support	151.00	80.00	0.74	231.00	231.00

6. Voting Ensembles for averaging the predictions for any arbitrary models.

```
In [52]: from sklearn.ensemble import VotingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
estimators = []
log_reg = LogisticRegression(solver='liblinear')
estimators.append(('Logistic', log_reg))

tree = DecisionTreeClassifier()
estimators.append(('Tree', tree))

svm_clf = SVC(gamma='scale')
estimators.append(('SVM', svm_clf))

voting = VotingClassifier(estimators=estimators)
voting.fit(X_train, y_train)

evaluate(voting, X_train, X_test, y_train, y_test)
```

Training Results:

Confusion Matrix:

```
[[327  22]
 [ 82 106]]
```

Accuracy:

0.8063

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	0.80	0.83	0.81	0.81	0.81
recall	0.94	0.56	0.81	0.75	0.81
f1-score	0.86	0.67	0.81	0.77	0.80
support	349.00	188.00	0.81	537.00	537.00

Testing results:

Confusion Matrix:

```
[[131  20]
 [ 38  42]]
```

Accuracy:

0.7489

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	0.78	0.68	0.75	0.73	0.74
recall	0.87	0.53	0.75	0.70	0.75
f1-score	0.82	0.59	0.75	0.71	0.74
support	151.00	80.00	0.75	231.00	231.00

In []: